

Divide and Conquer Networks

Joan Bruna and Alex Nowak

Courant Institute of Mathematical Sciences
Center for Data Science
NYU




NYU

**COURANT INSTITUTE OF
MATHEMATICAL SCIENCES**


Motivation

- We consider discrete geometric/algorithmic tasks:


Community Detection




Travelling Salesman




Shortest Paths



Sorting




Common Substrings



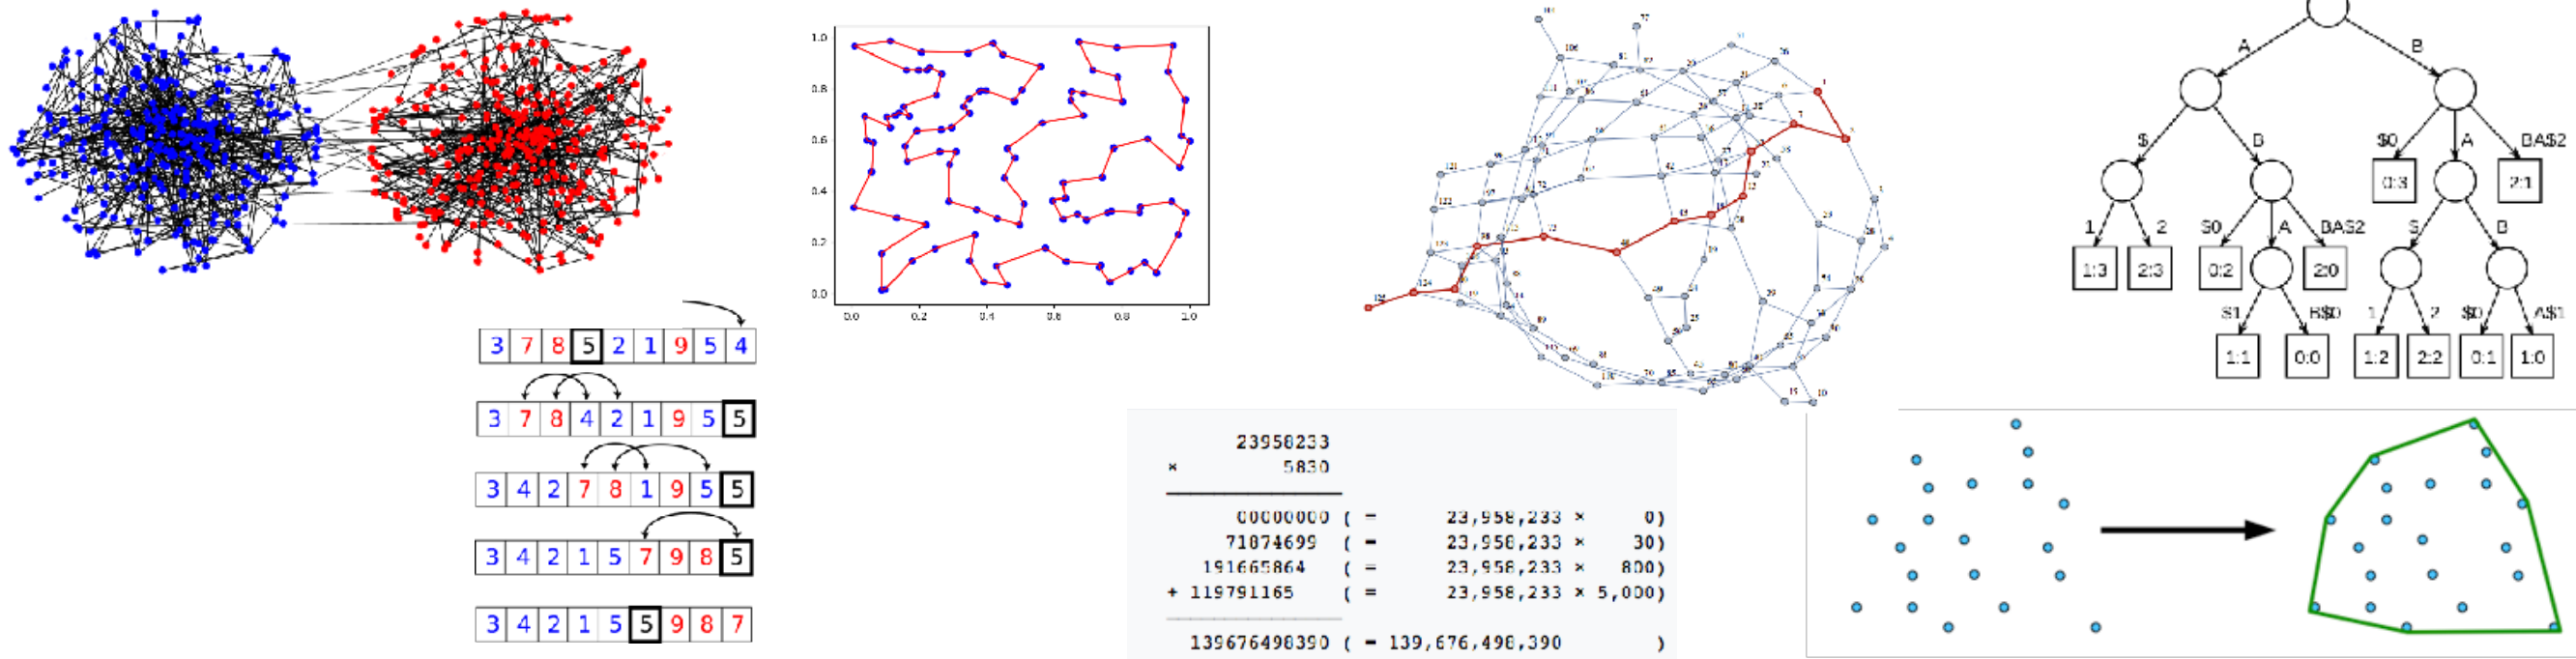
Multiplication

			23958233	
x			5830	
<hr/>				
	00000000	(=	23,958,233 x	0)
	71074699	(=	23,958,233 x	30)
	191665864	(=	23,958,233 x	800)
+	119791165	(=	23,958,233 x	5,000)
<hr/>				
	139676498390	(=	139,676,498,390)

Convex Hulls

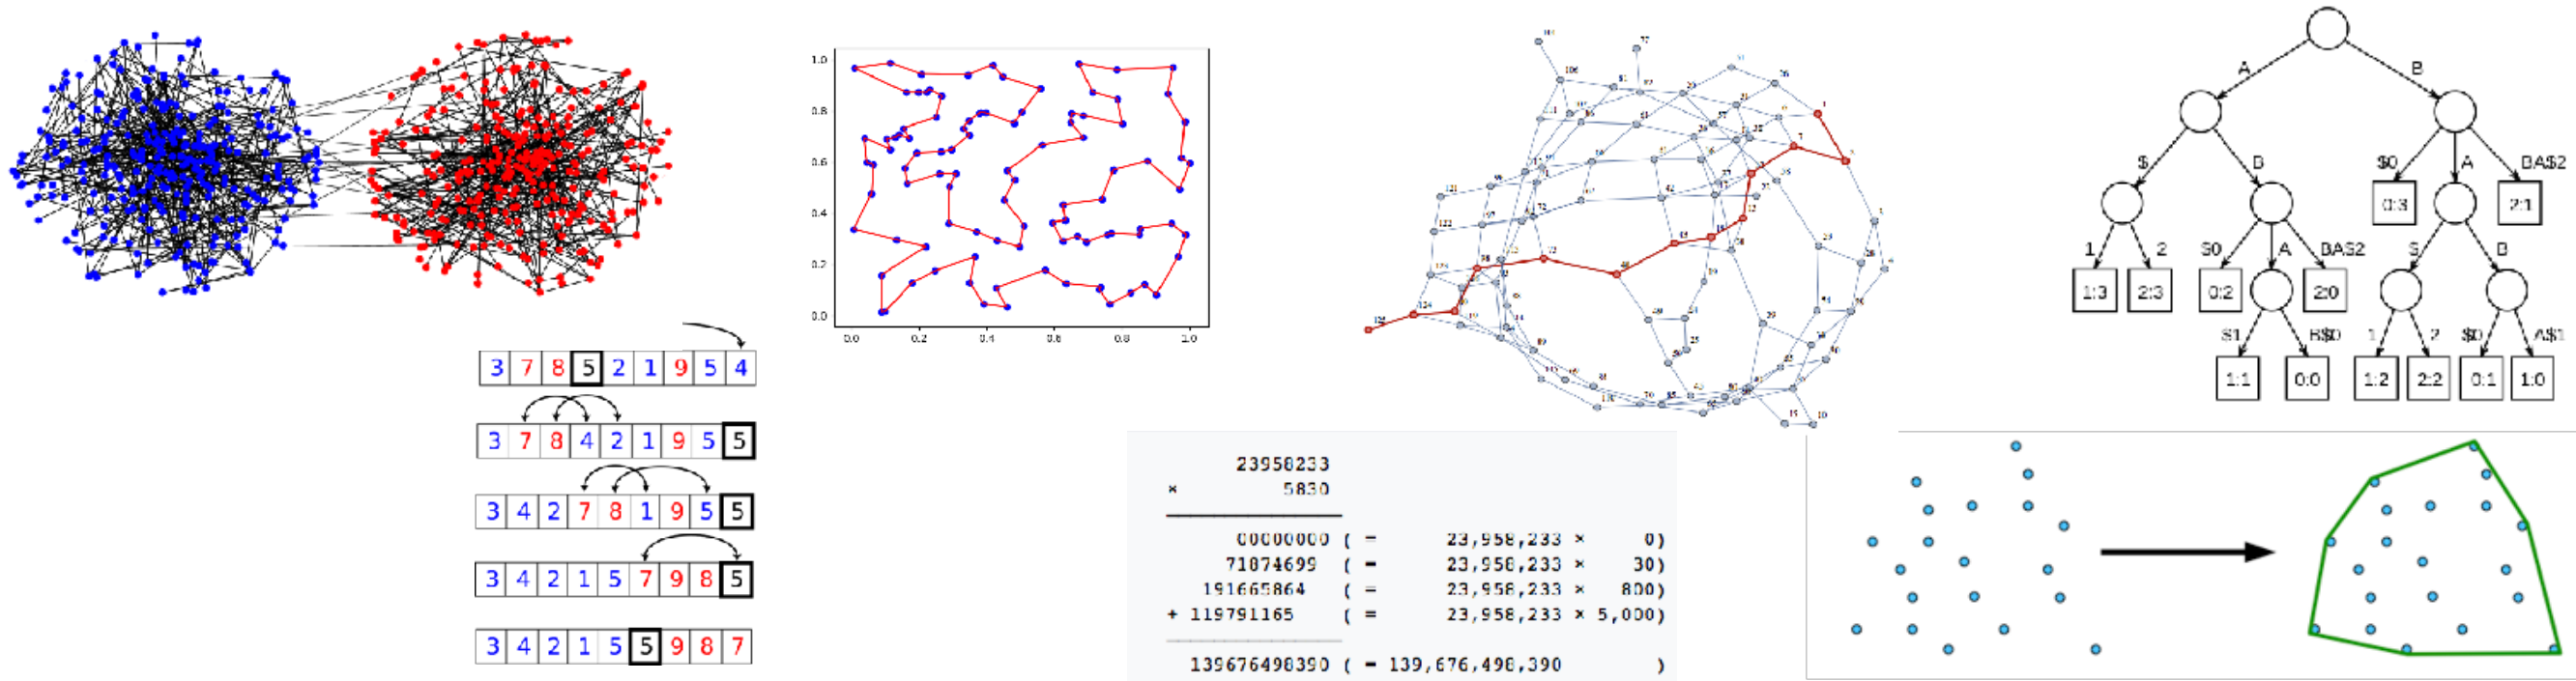


Motivation



- Tasks with known optimal worst-case complexity (e.g. sorting): can we leverage data distribution to obtain faster algorithms?
- Tasks with unknown optimal complexity (e.g. multiplication): can we learn how to obtain the fastest algorithm?
- NP-hard tasks (e.g. TSP): can we learn efficient approximations?

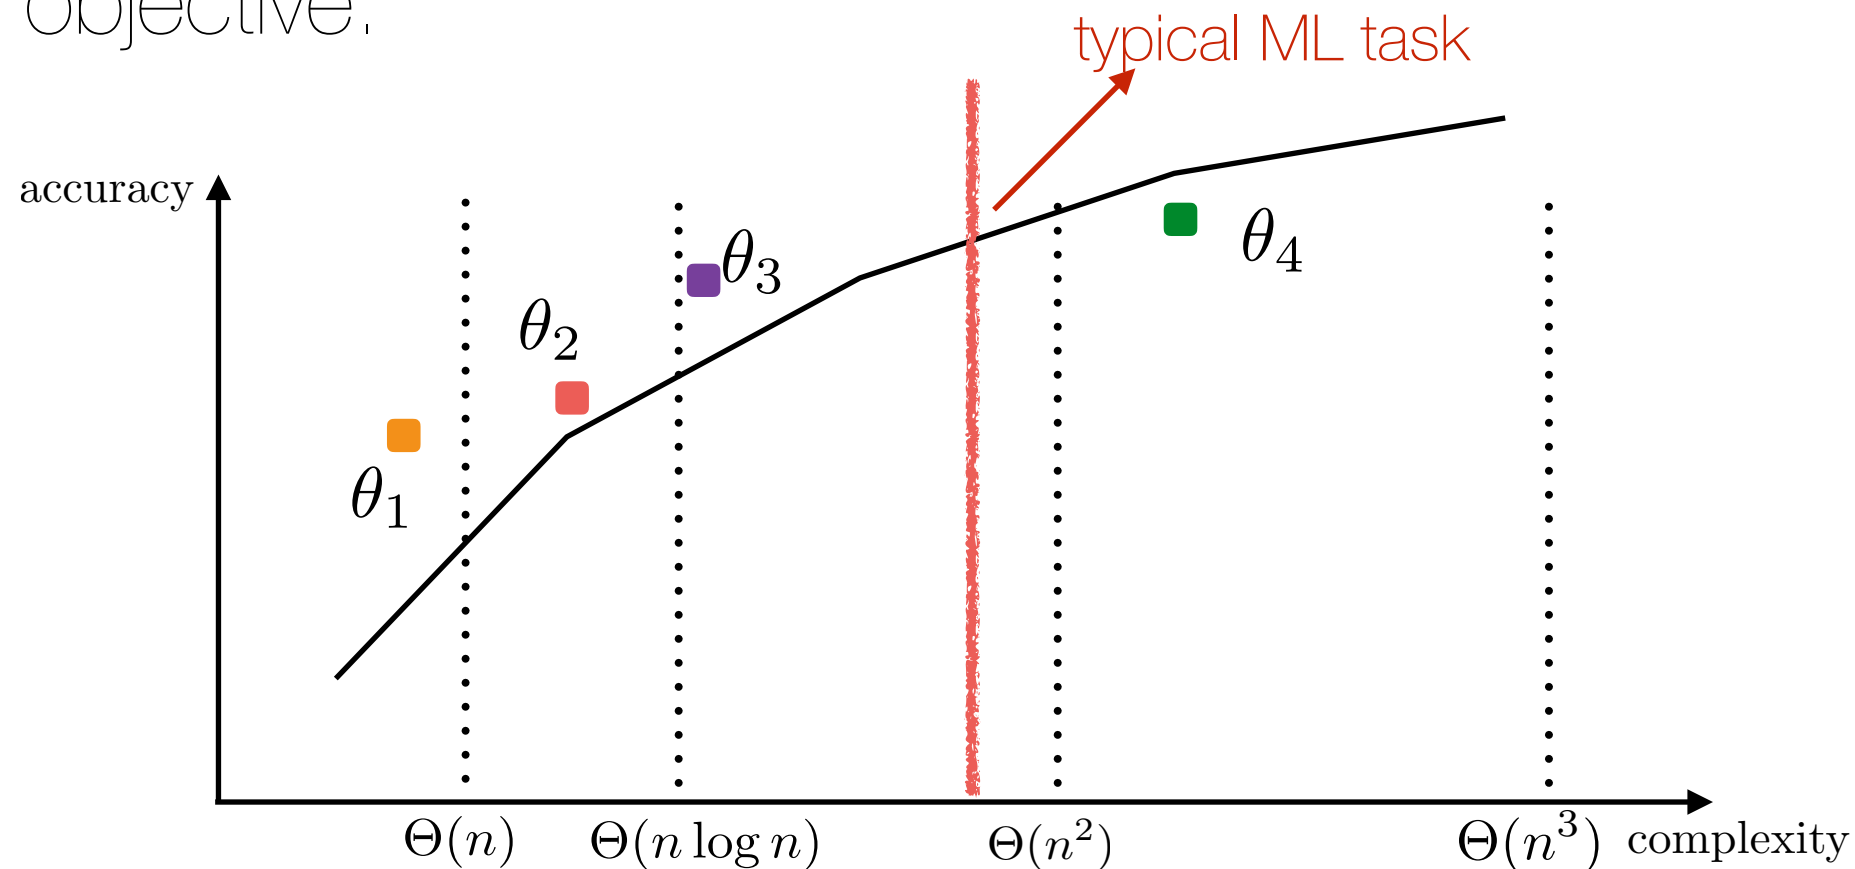
Motivation



- Tasks with known optimal worst-case complexity (e.g. sorting): can we leverage data distribution to obtain faster algorithms?
- Tasks with unknown optimal complexity (e.g. multiplication): can we learn how to obtain the fastest algorithm?
- NP-hard tasks (e.g. TSP): can we learn efficient approximations?
- **Goal:** learn how to solve these tasks *efficiently* from only input-output examples.
 - Need to generalize relative to input size.

Optimizing for Accuracy and Complexity

- We want to add the computational complexity as part of the learning objective.



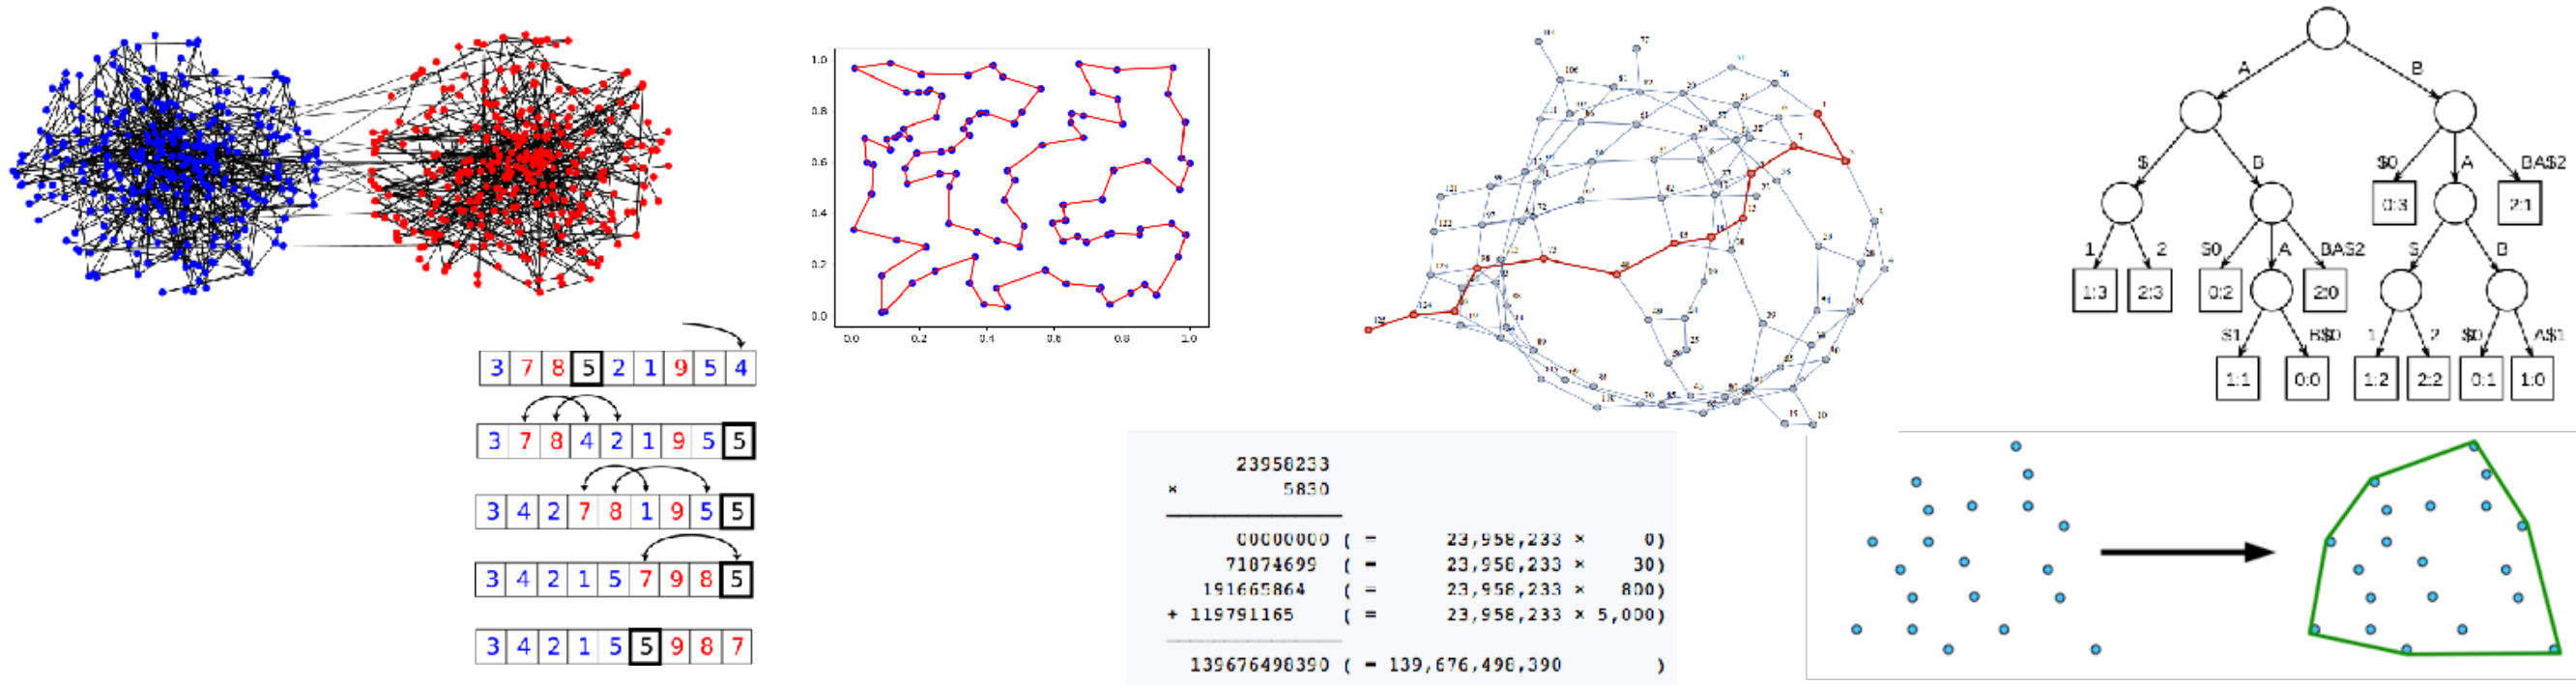
$$\mathcal{L}(\theta) = \mathcal{E}(\theta) + \lambda \mathcal{C}(\theta) .$$

enforces solving
the task with high
accuracy

enforces solving
the task with small
complexity

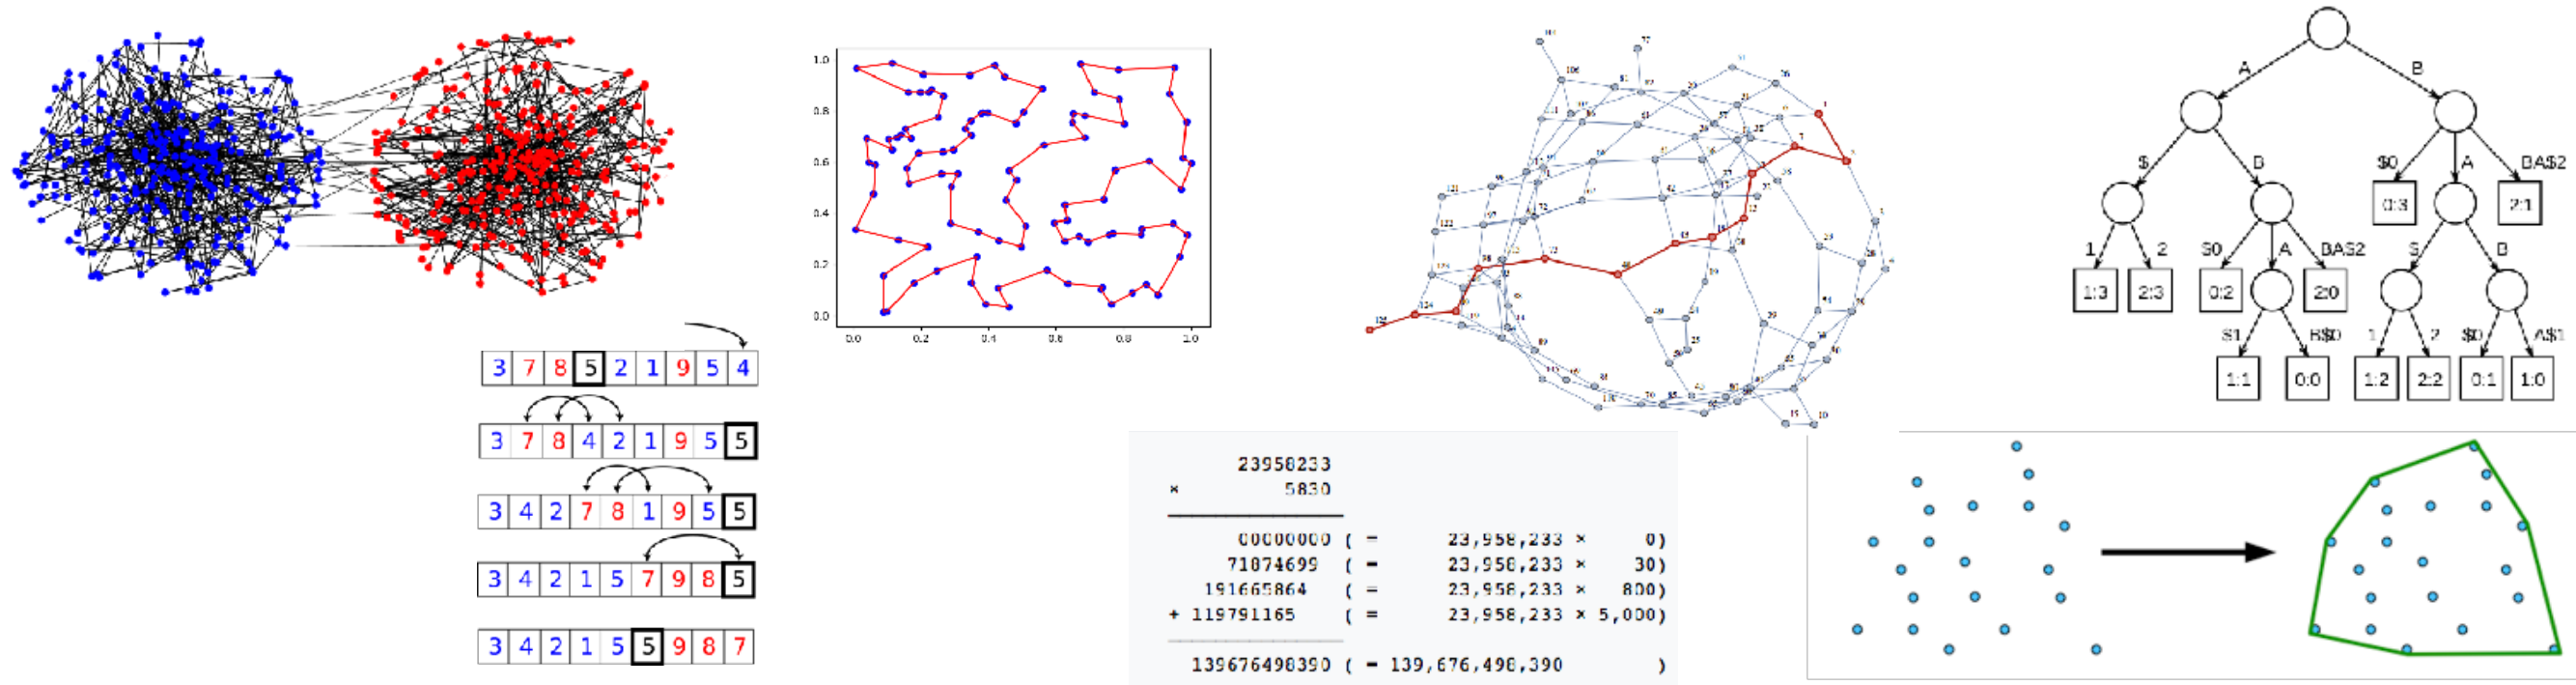
- Q: How to parametrize dynamic computations?

Scale Invariant Algorithmic Tasks



- Many algorithmic and geometric tasks are ***self-similar across scales***:
 - The solution can be expressed in terms of smaller input solutions.
 - This is the basis of dynamic programming and ***recursive*** algorithms.
 - Controlling the recursion = controlling computational complexity.

Scale Invariant Algorithmic Tasks



- Many algorithmic and geometric tasks are ***self-similar across scales***:
 - The solution can be expressed in terms of smaller input solutions.
 - This is the basis of dynamic programming and ***recursive*** algorithms.
 - Controlling the recursion = controlling computational complexity.
- We propose ***Divide and Conquer Networks***: a dynamic neural architecture that learns how to solve tasks using recursion.

Divide and Conquer Networks

- Consider general tasks \mathcal{T} that map an input X to output $\mathcal{T}(X)$.
input size $|X| = n$
- We consider a recursive decomposition of the task:

$$\mathcal{T}(X) = \mathcal{M}(\mathcal{T}(S_1(X)), \mathcal{T}(S_2(X))) ,$$

$$|S_i(X)| < n , \quad X = S_1(X) \cup S_2(X) .$$

Divide and Conquer Networks

- Consider general tasks \mathcal{T} that map an input X to output $\mathcal{T}(X)$.
input size $|X| = n$


- We consider a recursive decomposition of the task:

$$\mathcal{T}(X) = \mathcal{M}(\mathcal{T}(S_1(X)), \mathcal{T}(S_2(X))) ,$$


$$|S_i(X)| < n , \quad X = S_1(X) \cup S_2(X) .$$

- Rather than learning \mathcal{T} directly with e.g. a sequence-to-sequence model, we learn two basic operations:

- How to **split** a given input into two disjoint subsets:




- How to **merge** two partially solved tasks into a larger one:



Divide and Conquer Networks


- These blocks are applied recursively and dynamically:



- Each input generates a different “execution tree”.

Divide and Conquer Networks

- These blocks are applied recursively and dynamically:

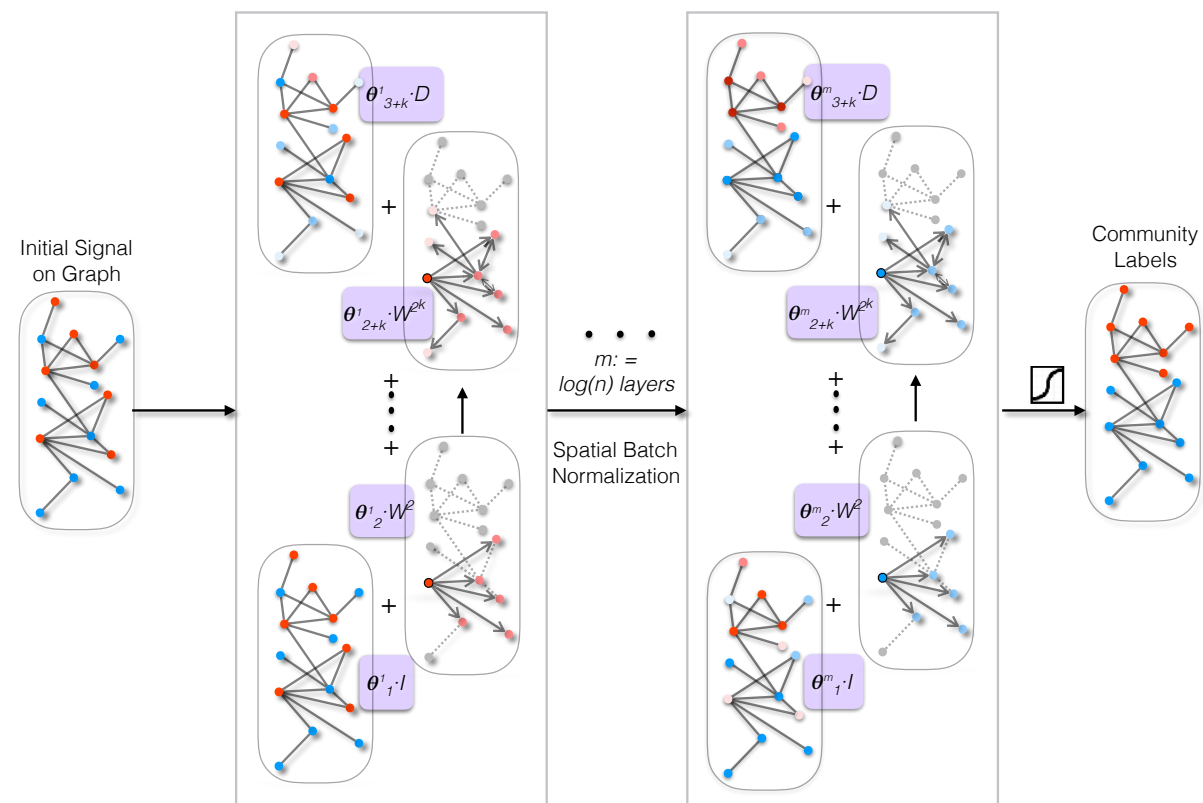


– Each input generates a different “execution tree”.

- Q: How to parametrize those operations?
- How to train the model end-to-end?

Split Model


- We consider split models that take as inputs either graphs or sets.
- For graphs, we consider Graph Neural Networks
 - First introduced in [Scarselli et al.'09], [Gori et al.'05].
 - Later simplified in [Li et al.'15], [Duvenaud et al.'15] [Sukhbaatar et al.'16].
 - Intimately related to convolutional neural networks using Graph Laplacians [Bruna et al.'16].



- For sets, we consider *set2set* modules [Vinyals et al.'16], [Sukhbaatar et al.'16]

Merge Model

- Since solutions may be partially ordered, we use a seq2seq model with attention [Bahdanau et al].
 - Generic block with $\Theta(n^2)$ complexity, but can be modulated by regularization.
- In the case where outputs are subsets of input (e.g. convex hull, TSP, sorting), the model is the so-called *Pointer-net* [Fortunato, Vinyals, '15].



- In that case, cascading merge blocks produces a product of stochastic permutation matrices at each scale.

Training the Model

- Training set: $\{(X^l, Y^l)\}_{l \leq L}$

$\mathcal{P}(X)$: partition tree associated to X

θ : split parameters

Accuracy loss:

ϕ : merge parameters

$$\mathcal{E}(\theta, \phi) = \frac{1}{L} \sum_{l \leq L} \mathbb{E}_{\mathcal{P}(X) \sim S_{\theta}(X)} \log p_{\phi}(Y^l \mid \mathcal{P}(X^l)) .$$

- Merge training:

– In a pointer network, $p(Y \mid X) = \Gamma X$, Γ : stochastic matrix.

– In our case,

$$p(Y \mid X) = \left(\prod_{j=0}^J \Gamma_j \right) X, \Gamma_j: \text{parameter sharing across scales.}$$

Training the Model

- Training set: $\{(X^l, Y^l)\}_{l \leq L}$
 $\mathcal{P}(X)$: partition tree associated to X θ : split parameters
Accuracy loss: ϕ : merge parameters

$$\mathcal{E}(\theta, \phi) = \frac{1}{L} \sum_{l \leq L} \mathbb{E}_{\mathcal{P}(X) \sim S_\theta(X)} \log p_\phi(Y^l \mid \mathcal{P}(X^l)) .$$

- Split training:
 - split parameters are separated from the output by sampling steps.
 - we use the policy gradient estimator:

$$\nabla_\theta \mathcal{E}(\theta, \phi) \approx \frac{1}{LK} \sum_{l \leq L} \sum_{\mathcal{P}(X^l)^{(k)} \sim S_\theta(X^l)} \log p_\phi(Y^l \mid \mathcal{P}(X^l)^{(k)}) \nabla_\phi \log p_\theta(\mathcal{P}(X^l)^{(k)} \mid X^l)$$

- variance is reduced thanks again to parameter sharing across scales.
- although other options (e.g relaxing sampling with softmax) possible too.

Computational Complexity as Regularization

- The average case complexity of splitting in our model is

$$C_S(n) = C_S(\alpha_S n) + C_S((1 - \alpha_S)n) + \Theta(n) ,$$

α_S : average fraction of elements sent to each split output.

- We verify that $C_S(n) \simeq \frac{n \log n}{\log \alpha_S^{-1}} .$

Computational Complexity as Regularization

- The average case complexity of splitting in our model is

$$C_S(n) = C_S(\alpha_S n) + C_S((1 - \alpha_S)n) + \Theta(n) ,$$

α_S : average fraction of elements sent to each split output.

- We verify that $C_S(n) \simeq \frac{n \log n}{\log \alpha_S^{-1}}$.
- Thus computational complexity is controlled by pushing α_S close to 0.5:

$$\mathcal{C}(\theta) = \left(n^{-1} \sum_{i \leq n} p_{\theta}(z_i \mid X) - \frac{1}{2} \right)^2 .$$

Experiments

- Sorting
- Convex Hull
- Clustering
- Community Detection on Networks
 - with L. Li (UC Berkeley)
- Quadratic Assignment Problem and TSP
 - with A. Nowak, S. Villar and A. Bandeira (NYU).

Sorting Real Numbers



- Toy-task to assess capacity of models.
- Baseline: Pointer Network applied directly to the input.
- DCN (no split). Fix split block to generate a balanced binary tree independent of X .
- DCN (no merge): Fix Merge block to the identity.
- DCN (Joint): Train both blocks.
- We use weak supervision and computational regularization. Train for $n=8, 16$. Test results:

	Baseline	DCN (no Split)	DCN (no Merge)	DCN (Joint)
$n=8$	80.1	90	100	100
$n=16$	31	67	100	100
$n=64$	0	0	99	0

– Joint model does not generalize because merge block does not have the appropriate complexity for this task ($\Theta(n^2)$ vs $\Theta(n \log n)$).

Convex Hull

- Given n points in \mathbb{R}^d , find the extremal set of points of the polytope of minimum area that contains them all.
 - Task can be solved efficiently in $O(n \log n)$ in the plane




	n=25	n=50	n=100	n=200
Baseline	81.3	65.6	41.5	13.5
DCN Random Split	59.8	37.0	23.5	10.29
DCN	88.1	83.7	73.7	52.0
DCN + Split Reg	89.8	87.0	80.0	67.2


Baseline: seq-2-seq model (PtrNet)

Convex Hull

- 2 Scales (Random Split):



- 2 Scales (Learnt Split):



Euclidean K-Means

- Given n points of \mathbb{R}^d , solve the following combinatorial optimization problem:

$$\min_{\mathcal{P}(X)} \sum_{i \in \mathcal{P}(X)} n_i \sigma_i^2$$

σ_i^2 : variance of subset i of partition projected on an (unknown) subspace.
 n_i : cardinality of subset i .

Euclidean K-Means




- Given n points of \mathbb{R}^d , solve the following combinatorial optimization problem:

$$\min_{\mathcal{P}(X)} \sum_{i \in \mathcal{P}(X)} n_i \sigma_i^2$$

σ_i^2 : variance of subset i of partition projected on an (unknown) subspace.
 n_i : cardinality of subset i .

– Here the output is only specified with split parameters.


	k=4, n=40	k=8, n=80	k=16, n=160
Baseline	0.35	4.62	16.08
DCN	0.85	0.86	0.86




baseline: train a direct policy ($J=1$)

Euclidean K-Means

- Baseline, $k=8$ and $n=80$.



- DCN, $k=8$ and $n=80$.



Summary So Far

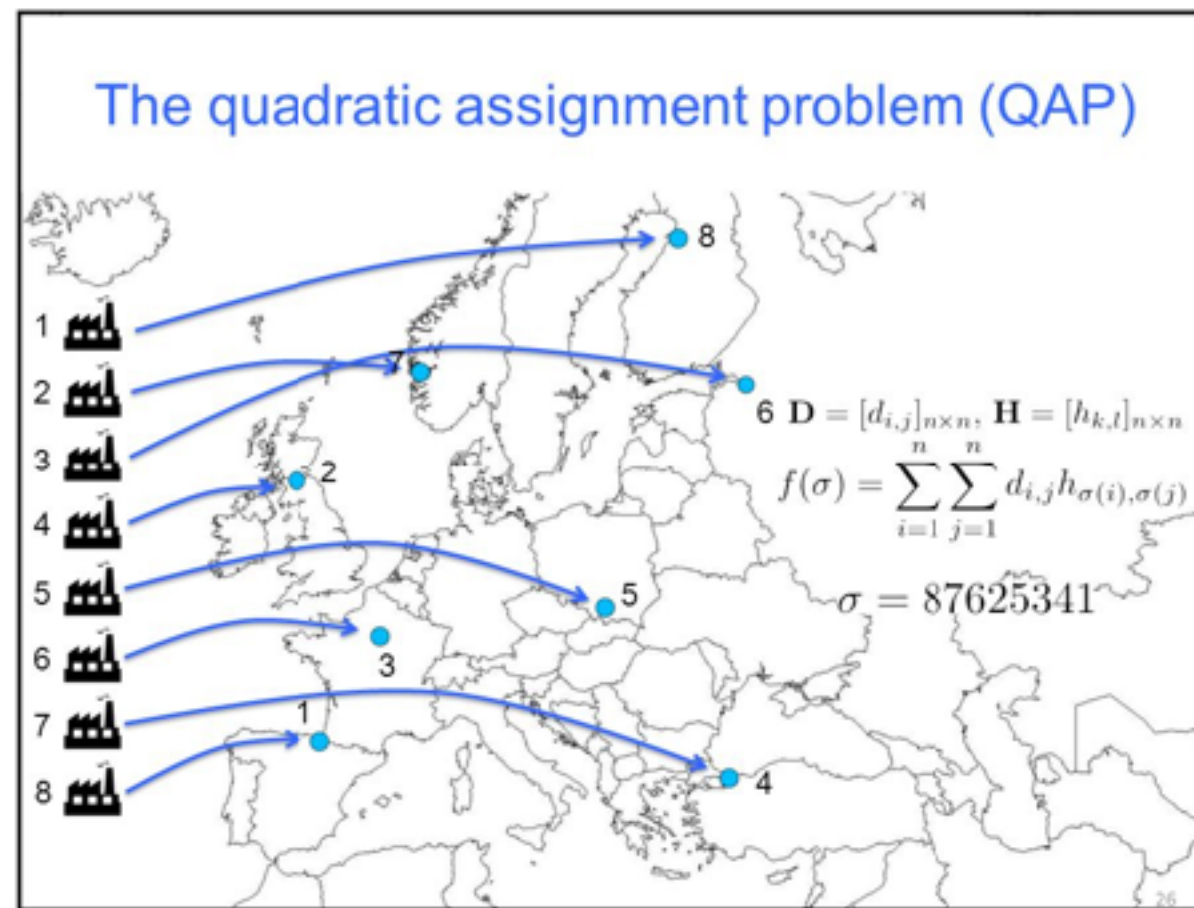
- So far, we have shown a general trainable architecture to exploit scale invariance in algorithmic/ geometric tasks.
 - Significant gains over baselines agonistic to recursion.
 - Computational and Accuracy Tradeoffs by gradient descent.
 - Shown on tasks where we currently have optimal, efficient algorithmic solutions.

Summary So Far

- So far, we have shown a general trainable architecture to exploit scale invariance in algorithmic/ geometric tasks.
 - Significant gains over baselines agonistic to recursion.
 - Computational and Accuracy Tradeoffs by gradient descent.
 - Shown on tasks where we currently have optimal, efficient algorithmic solutions.
- Current work: use this machinery on tasks that are computationally and statistically *hard*, or where no optimal algorithm is known.
 - Community Detection in Noisy Graphs.
 - Quadratic Assignment Problem
 - ❖ Travelling Salesman Problem
 - Matrix Multiplication.

Quadratic Assignment Problem

[with S. Villar (NYU), A. Nowak (NYU) and A. Bandeira (NYU)]



- Find an assignment that optimizes the transportation cost between two graphs:

$$\min_{X \in \Pi_n} \text{Tr}(A_1 X A_2 X^T) .$$

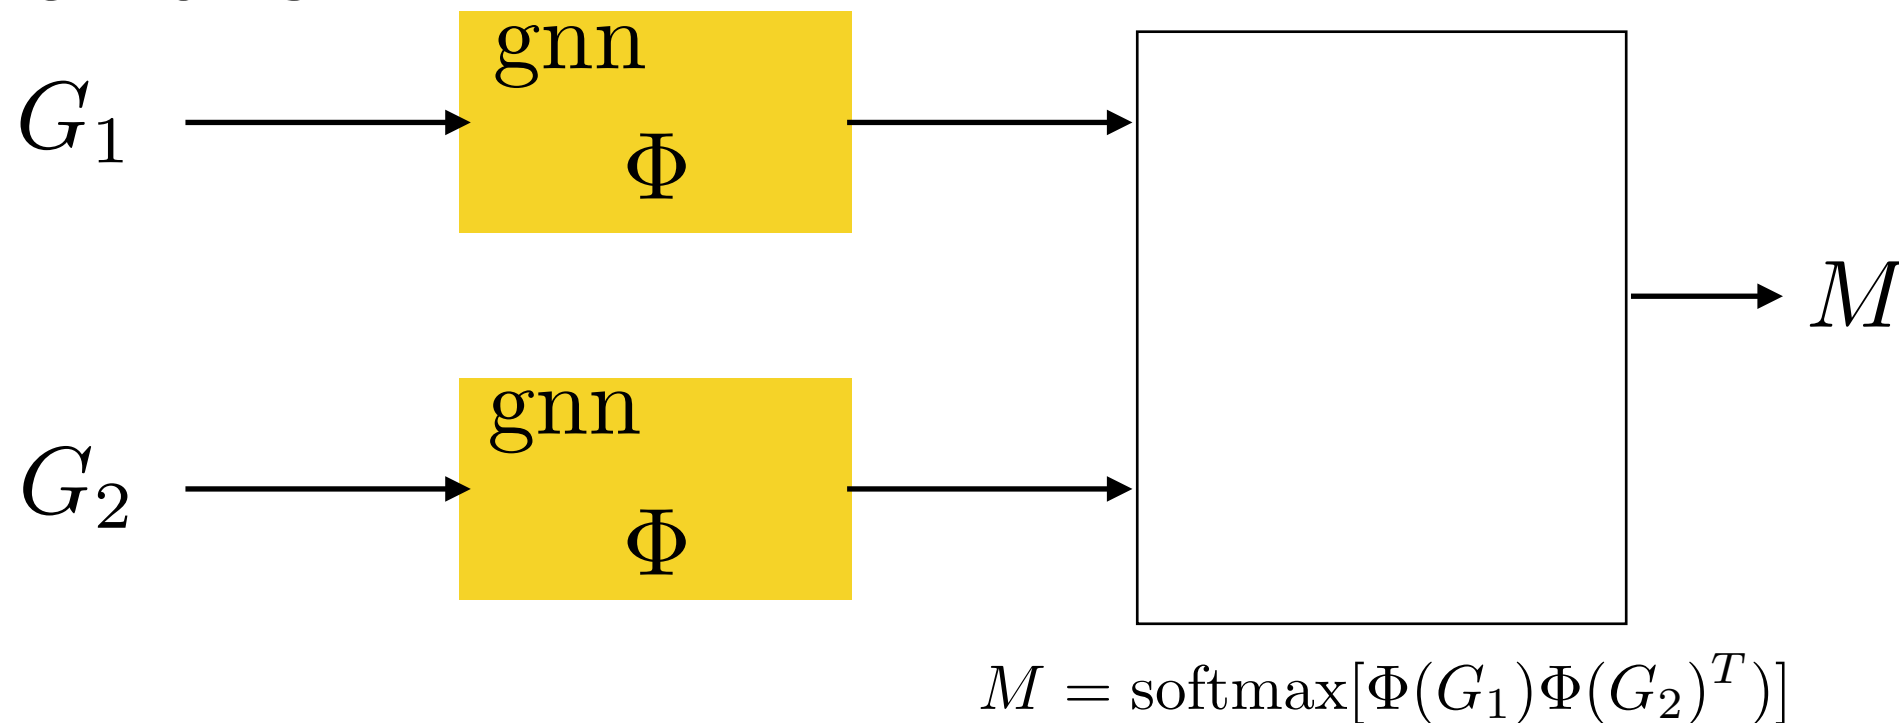
Π_n : space of $n \times n$ permutation matrices.

- NP-hard
- Contains the TSP as a particular instance.
- Relaxations using SDP and Spectral Approaches.

Quadratic Assignment Problem

[with S. Villar (NYU), A. Nowak (NYU) and A. Bandeira (NYU)]

- We learn approximate solutions using siamese graph neural networks:



- We train the model to predict the correct permutation matrix on a dataset of the form

$$G_1 = PG_2 + N$$

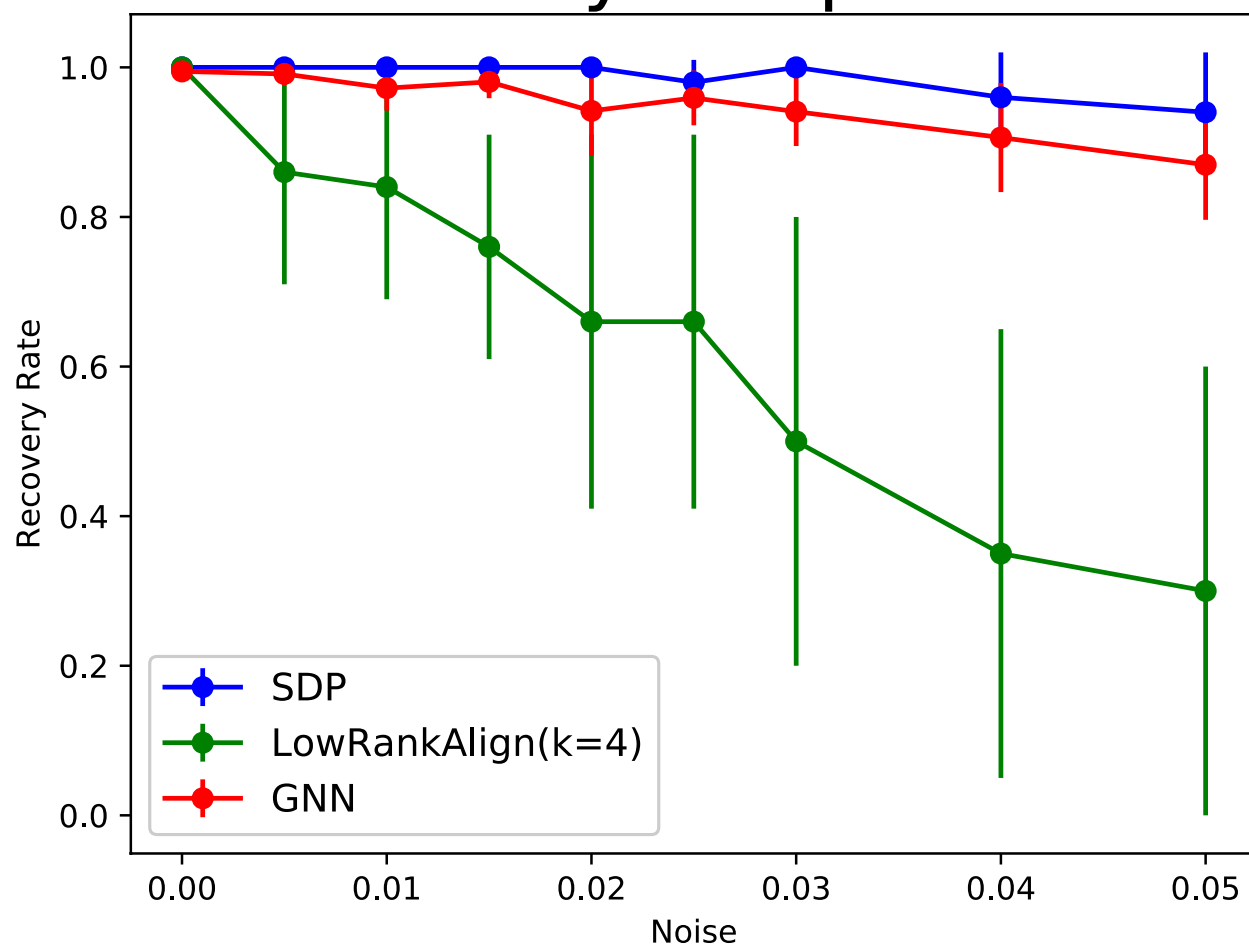
$$N \sim \text{Erdos-Renyi}$$

- $G_2 \sim \text{Erdos-Renyi}$
- $G_2 \sim \text{Random Regular}$


Quadratic Assignment Problem

[with S. Villar (NYU), A. Nowak (NYU) and A. Bandeira (NYU)]

ErdosRenyi Graph Model



Random Regular Graph Model



- Our model runs in $o(n^2)$
- LowRankAlign is $o(n^3)$
- SDP runs in $o(n^4)$

Travelling Salesman Problem

[with S. Villar (NYU), A. Nowak (NYU) and A. Bandeira (NYU)]

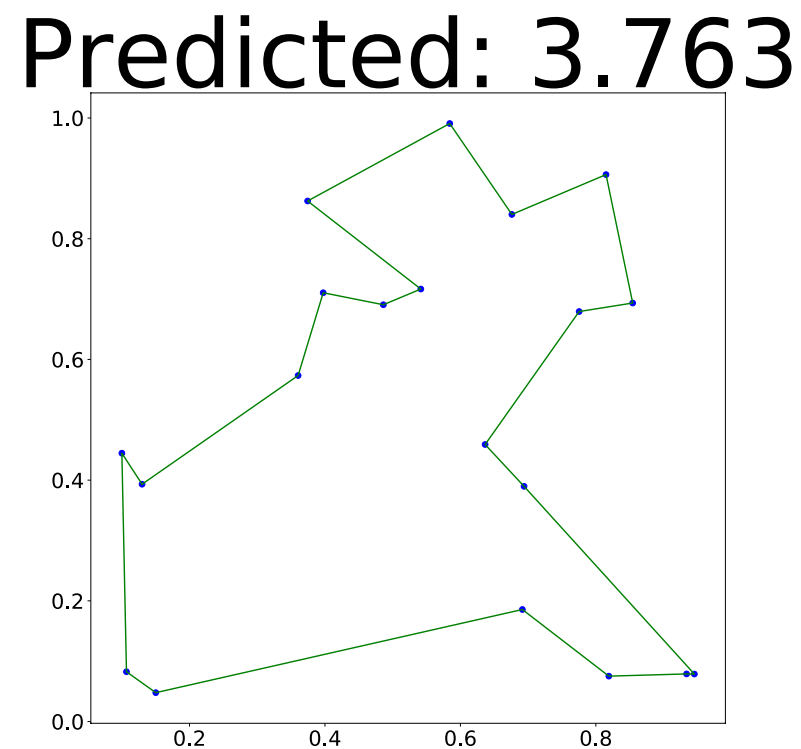
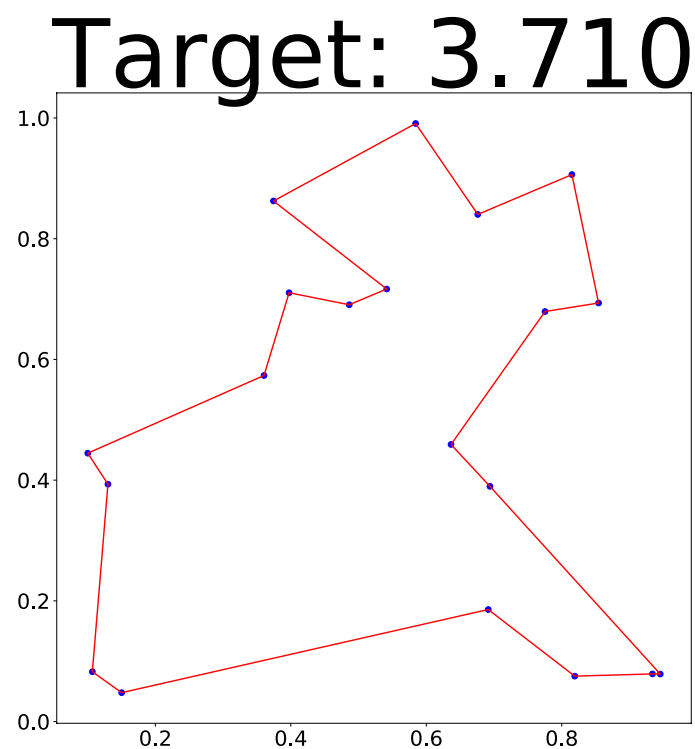
- A particular case of the QAP is the Travelling salesman Problem:

$$\min_{X \in \Pi_n} \text{Tr}(A_1 X A_2 X^T) \text{ .}$$

A_1 : dense pairwise distance matrix

A_2 : n -cycle adjacency matrix

- We obtain preliminary good results



- Not yet like best Heuristics in the planar case (Christofides algorithm).
- Ongoing work.

Conclusions

- Neural Networks and Dynamic Programming.
 - Many tasks in real life contain some form of scale invariance (e.g. navigation, planning, scheduling, ranking, ...).
 - Real need for accuracy-complexity tradeoffs.
 - Our architecture is a step towards end-to-end learning that exploits such structure.
 - Challenges ahead: more flexible models allowing wider range of accuracy-complexity tradeoffs.
- Neural Networks and Complexity.
 - Surprising, Intriguing ability to reach detection thresholds known to be hard on the community detection.
 - Very efficient data-driven alternative to attack hard combinatorial optimization problems.
 - Challenges ahead: understand what are the fundamental limits of data-driven

Thank you!

References:

“Divide and Conquer Networks”, A. Nowak and J. Bruna, submitted. (<https://arxiv.org/abs/1611.02401>)

“Community Detection with Graph Neural Networks”, J. Bruna and L. Li (<https://arxiv.org/abs/1705.08415>)

“A Note on Learning Algorithms for Quadratic Assignment with Graph Neural Networks”, A. Nowak, S. Villar, A. Bandeira, J. Bruna, submitted.